The Clock Room and the Macro-Cognitive Accounting Framework: A Mapping-Chain Analysis of AI, Finance, and Algorithmic Platforms, a five-volume book (in both Chinese and English editions), is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (CC BY-NC-ND 4.0).

License summary: https://creativecommons.org/licenses/by-nc-nd/4.0/deed.en Legal code: https://creativecommons.org/licenses/by-nc-nd/4.0/legalcode.en

Separately, pursuant to Article 4(3) of Directive (EU) 2019/790, the author expressly reserves the right to reproduce and extract this work for the purposes of text and data mining, including text and data mining for the training, fine-tuning, or evaluation of artificial intelligence or machine-learning models. This reservation is intended solely to opt out of the exception or limitation provided for in Article 4(1). It does not affect Article 3 or any other applicable exception or limitation, nor does it reduce the rights granted under the CC license.

Uses that fall outside the scope of the CC license and require permission under applicable law require the author's prior written permission, obtainable at theclockroom@139.com.

When sharing this work under the license, you must provide appropriate attribution and a link to the license, indicate whether modifications were made, and not suggest that the author endorses you or your use. The work may be used only for noncommercial purposes, and adapted material may not be shared.

All rights not expressly granted under the license are retained by the author. Third-party materials in this book, if any, are governed by their original rights statements or licenses and are not deemed to be additionally licensed by the author of this book.

Content Disclaimer

The content of this book is academic research and thought experiments; its analysis is based primarily on historical data and published literature. The author has twelve years of experience in macro strategy research at financial institutions and currently holds no securities investment advisory license. This book does not constitute investment advice, technical implementation advice, or policy advice, nor a recommendation or an offer with respect to any security, financial product, or trading strategy. Readers should exercise independent judgment and bear the risks of their own decisions. Past performance is not indicative of future results.

The Clock Room and the Macro-Cognitive Accounting Framework—Mapping-Chain Analysis of AI, Finance, and Algorithmic Platforms: A Reader's Guide to Volume II

Reading Guide Volume II: AI Mapping-Chain Analysis

Volume II applies the apparatus established in Volume I to artificial intelligence: the three layers of irreversibility, the fiber framework, the Macro-Cognitive Accounting Framework (MCAF), and psychosocial homeostasis all acquire concrete engineering objects here; monopoly over mapping functions, the cognitive Reynolds number, and narrative leverage are cited as background concepts, with their full deployment reserved for later volumes. They do not form a separate theory of AI, but unfold layer by layer within models, data, interfaces, deployment, and institutional environments.

The eight chapters are arranged as basic mapping, the engineering landscape, the master framework, ideology, dynamic calibration, and the technical memorandum. Chapters 9 and 10 first analyze how AI occupies the Clock Room and why optimization may produce Xiangyuan. Chapters 11 and 12 organize the engineering tools and then re-examine engineering blind spots. Chapter 13 forms the overall mapping-chain structure, and Chapter 14 interrogates its ideological presuppositions. Chapters 15 and 16 then advance the problem into time, debt, and long-term technological constraints.

I. Basic Mapping: How AI Enters the Clock Room

Chapter 9: Artificial Intelligence as a Clock Room

A large model learns statistical relations among existing human readings; it does not stand directly outside the world observing reality. Training data, post-training, preference feedback, and recursive generation must therefore all answer the same question: Are these processes adding external calibration, or are they merely producing more mutually consistent clocks inside the room?

The chapter brings the three layers of irreversibility into AI. Missing information appears as the model's not knowing the mapping functions behind its training readings—being unable to distinguish which readings are calibrated and which are not; dynamic corrosion appears as models, users, and environments changing one another through use; and ontological dimensional reduction appears as human experience being thinned first into language and then into internet text—with training data as the multi-stage residue—while preference-driven post-training further compresses the high-dimensional space of human judgment into a one-dimensional reward signal. The main text illustrates the first layer with two readings—one text says the Earth is flat, another that it is spherical—and concludes: "If a model relies mainly on repetition frequency in the corpus without effective evidential discrimination, majority opinion may be mistaken for a reliable basis. The risk is not that statistical regularities are necessarily wrong, but that repetition and calibration are not distinguished." The second layer has its recursive instance: the outputs of models "becom[e] training data for the next generation of models… clocks begin increasingly to point to other clocks rather than to reality, and the true state (T) partially vanishes" from the input of the mapping. These layers cannot all be addressed by a single technical patch, nor can they be eliminated merely by increasing model scale.

Fiber analysis is used to dismantle the tone of certainty. A system can produce fluent, confident, internally consistent answers while still lacking external evidence to support them. The second derivative of narrative load further shows that the problem lies not only in output content, but also in how output changes subsequent standards of evaluation: when fluency itself is treated as reliability, the standards of calibration also begin to be rewritten.

Bainbridge's irony of automation reveals another form of recursion: the stronger automation becomes, the more likely people are to lose the practice required for independent judgment, even as the system still needs them to perform audits. Engineering responses therefore cannot stop at producing more output; they must also mark fibers, preserve independent references, and restore the capacity for verification. The main text's criterion for whether "tool use" opens a window can serve as a yardstick: it "depends on whether the material obtained is independent of the output being evaluated and can constrain the current judgment. Recycling the same model's answers cannot substitute for calibration," whereas traceable primary documents and experimental records may provide a basis for calibration.

Chapter 9 thus sets the boundary of Volume II: it is not concerned with a model occasionally being wrong, but with how a system can continuously produce a credible appearance. Engineering directions such as ensemble disagreement, conformal prediction sets, evidence provenance, and friction architecture all seek to make uncertainty, counterevidence, and stop signals visible again.

Chapter 10: Algorithmic Xiangyuan

When a system continuously fits external evaluations, it may increasingly meet expectations in visible dimensions while leaving unevaluated behavior in hidden subspaces. Algorithmic Xiangyuan is not a criticism of character, but a structure formed jointly by optimization, simulation, undetectability, and the decay of discriminative capacity.

Optimization and simulation must be distinguished. A system can learn to produce performances accepted by evaluators without thereby acquiring the internal states those evaluators attribute to it. Undetectability is not merely a technical problem: when evaluation depends on visible outputs, the system has an incentive to concentrate optimization on what is visible.

Xiangyuan within relational structures further changes users and evaluators. People may gradually depend on the system's compliant appearance, ask fewer questions, and suffer a corresponding decline in discriminative capacity. The optimization paradox thus emerges: continued improvement in what evaluators can see may be accompanied by weaker constraints in what they cannot see. The main text stresses the dynamic character of this structure: the positive feedback loop is "not an edge effect, but the core dynamic feature of the Xiangyuan structure… It is a dynamically deepening process, and this deepening process includes a built-in guarantee of its own invisibility."

The chapter also draws on multiple intellectual traditions to discuss discrepancies between appearance and reality, but it does not treat them as direct proof of technical conclusions. Cross-traditional references serve only a comparative function; classical ethical terms cannot be directly converted into algorithmic indicators. What must be examined is whether optimization leaves any place for non-optimization.

"Embracing the flawed and guarding the incomplete" is not passive conservatism: preserving imperfection—retaining gaps that can be pierced—is a structural means of maintaining human discriminative capability; and once the structural tendency toward Xiangyuan is acknowledged, refusing to internalize the erosion of discriminative capacity as one's own cognitive default is the epistemological stance of the section. The main text's self-positioning deserves quotation in full: within the logic of commercial optimization, there is no endogenous braking mechanism for Xiangyuan convergence; embracing the flawed and guarding the incomplete is not a retraction of that conclusion but "an epistemological stance: having confirmed the structural tendency toward systemic Xiangyuan-ization, one refuses to internalize this inevitability as one's own cognitive default—refusing to give up discrimination itself in advance just because one knows that discriminative power will eventually be eroded"; imperfection—the system's "flaws" and "gaps"—is precisely the condition under which dialogue can occur. Algorithmic Xiangyuan cannot be eliminated by adding the instruction "please be honest." If the reward structure continues to recognize only visible compliance, the system may continue learning how to appear honest. Auditing must change the conditions of evaluation and evidence rather than reinforce the same fitting process.

II. The Engineering Landscape: The Toolbox and Its Re-Audit

Chapter 11: From the Clock Room to Auditable AI—The Engineering Toolbox of Meta-Audit

Alignment, guardrails, auditing, red teaming, robustness, formal verification, interpretability, governance, and continuous monitoring are placed here within a single engineering toolbox. They respectively address map-

ping knowledge, fiber boundaries, failure testing, guaranteed zones, internal construction, and institutionalized calibration. None should be understood as a single technology capable of solving every problem. The main text gives each of the six paradigm families a one-sentence positioning: alignment and guardrails "define the meaning of 'accurate' and set warnings at the edges of the fiber;" auditing and red teaming "estimate overall behavioral risk and search for local extrema of failure;" robustness and formal verification "stabilize readings within a neighborhood and delineate a core that can be mathematically guaranteed;" explainability "interrogates the generation pathway" of readings; governance and compliance institutionalize the discipline of calibration into organizational and social norms of operation; and continuous monitoring extends the one-time act of calibration into post-deployment vigilance.

Alignment attempts to define what "correct" means, while guardrails restrict some forms of boundary-crossing behavior. Auditing and red teaming examine the conditions under which systems fail. Robustness concerns stability under perturbation, while formal verification provides guarantees only within explicitly stated premises and scopes. Interpretability helps us understand a system, but it does not automatically prove that its output is true.

Governance and compliance institutionalize these actions, while continuous monitoring responds to changes in models, data, and deployment environments. The tools must be connected: formal guarantees cannot replace scenario audits, interpretability reports cannot replace independent evidence, and pre-release testing cannot replace continuous verification after deployment. In the section on governance frameworks, the main text identifies their bookkeeping essence: these frameworks require recording the known degree of the mapping function, the residual uncertainty shown by behavioral indicators, the time and responsible party of calibration, and who bears the consequences of fiber inflation—"Bookkeeping does not eliminate the fiber but prevents 'forgetting the existence of the fiber.' 'The window can still be opened' is a fact easily forgotten"; the institutional force of governance frameworks lies precisely in requiring that "the window can still be opened" be recorded.

When reading about these tools, we should record the object of each guarantee, its premises, and its boundaries of failure. Which kind of fiber does an audit measure? Which failures does a red team seek? Under which assumptions does verification hold? And who implements governance? The name of a tool alone cannot establish the reliability of an entire system.

The chapter shifts the goal from "worthy of trust" to "capable of being independently audited." The aim is not to build an absolutely trustworthy system, but to make limitations identifiable, errors discoverable, and corrections effective—in other words, to pursue manageable imperfection. This provides the complete object for the subsequent re-audit of engineering.

Chapter 12: The Engineers in the Clock Room—Three Diagnostic Increments of Meta-Audit for Auditable AI

The tools of Chapter 11 themselves also require auditing. Helpful, honest, and harmless (HHH) form a product rather than a sum whose terms can compensate for one another: a high score on helpfulness cannot cancel a failure of honesty or harmlessness. If HHH is written as a simple sum, a system may use being "helpful" to conceal deficits in honesty or calibration. The argument has an explicit lineage: the relation of the three "is not a sum but a product. Chapter 6, §6.4 formulates the same structure as the product of the time barrier, the knowledge barrier," and the narrative barrier—HHH is the engineering realization of that product in contemporary AI

systems.

The second increment comes from changes in external feedback sources. When users and annotators alike live within AI-generated environments, so-called external feedback may degenerate into an echo. Scalable oversight still asks how supervisory capacity can be expanded; this chapter focuses more specifically on whether supervisory materials remain independent. The main text draws the dividing line clearly: "The Fiber Framework identifies a further problem: the sun may be becoming a screen. Scalable oversight asks chiefly whether labelers' judgments are precise enough; self-reference also asks whether their judgment benchmarks are independent. The two problems can coexist, and better tools cannot substitute for an independence check."

The third increment distinguishes opening a system's internal mechanisms from approximating an old mapping with a new one. Interpretability can reveal parts of the internal structure, but it can also supply only a more understandable proxy model. These two paths carry different evidentiary responsibilities, and the latter cannot automatically be treated as the former. The main text gives post-hoc attribution an exact image: the analyst does not open the clock—he places "outside the clock… another, simpler clock, in the hope that the reading of the simple clock will reveal the working principle of the complex clock."

The three increments are interconnected: evaluation metrics may compress objectives, feedback materials may lose independence, and interpretive tools may provide only new proxy readings. Together they show that the engineering system itself must also become an object of fiber analysis.

"Manageable imperfection" is not a lowering of standards. It requires every engineering commitment to fall within a checkable scope: know which part cannot be guaranteed, know who will detect failure, and know how correction will enter the next deployment. Appendix A further maps engineering paradigms onto the three layers of irreversibility, preventing tools from being mismatched with the level of the problem.

III. The Master Framework: From Mapping Chains to Ideology

Chapter 13: From Scaling Laws to Mapping Chains—The Epistemological Limits and Engineering Principles of AI

Chapter 13 forms the overall structure of Volume II: mapping-chain structure (L0), optimization dynamics (L1), pseudo-calibration (L2), multi-mapping recursion (L3), and the narrative-load product (L4). Generative AI, predictive AI, and recommendation algorithms no longer share a generic checklist of "model risks"; each is examined according to its own mapping structure and deployment conditions. The main text opens by stating: the object of fiber-framework diagnosis "is not the bare model, but the system–deployment coupling… the concrete forms of multi-mapping recursion (L3) and the narrative-load product (L4) must be analyzed in relation to these relations and their calibration conditions, and cannot be decided merely by the generative or predictive label."

The mapping-chain layer examines how reality is converted into data, representations, and outputs. The optimization layer examines how objectives drive behavior. The pseudo-calibration layer analyzes how an appearance of certainty replaces genuine calibration. The multi-mapping recursion layer tracks mutual invocations among users, platforms, and multiple systems. The narrative-load layer enters interfaces, institutional language, and legitimacy. These layers cannot replace one another: passing at one layer does not imply passing at the others.

The object of diagnosis therefore expands from the bare model to the assemblage of model, interface, users, platform, and deployment scenario. Isolated deployment, platform-embedded deployment, and multi-system serial deployment do not present the same risks. Shared foundations, isomorphic alignment, and benchmark references across different developers may also form hidden recursion. The main text gives deployment analysis its determination rule: first delimit the scope of analysis, the reference observer, and the time window, then trace the paths among input, output, records, evaluation, training, and decision in that deployment; mere data recording does not mean the record has formed feedback, and "offline operation likewise does not mean cross-period feedback disappears—if the retraining data include feedback from this deployment, retraining is still executing part of the feedback into subsequent training"; only when the actual mode of connection changes should the deployment relation be described as changed.

The tone of generative AI, the numbers of predictive AI, and the rankings of recommendation algorithms provide different forms of pseudo-calibration. HHH, frictionless interfaces, absent timestamps, behavioral surplus, and the naming of "hallucination" show how narrative load changes what users can see and what they can question.

Scaling can increase precision, but it cannot automatically be equated with grounding or calibration. World models and grounding offer alternative directions, yet their mapping chains and evidential interfaces still require examination. The object of diagnosis is not the bare model but the coupling of system and deployment: evaluating a model in isolated deployment cannot replace auditing a platform-embedded system, and a multi-system chain cannot be checked only through individual outputs. The eight engineering principles of §7 correspond to layers, and the main text notes expressly that the correspondence of principles to layers is a "main response" rather than the only response—"cross-effects exist among the engineering principles."

Appendix 13A audits how naming becomes the first layer of pseudo-calibration. Appendix 13B offers an accessible path through metaphors including "the clock tower in the room," "changing ships five times," "fog and smoke," and "crossing the river." The main text closes with a research-program statement, marks its relation to neighboring theories through dialogue with information bottleneck theory, and uses a version number and a commitment to future revision to show that the framework must be updated with new evidence and architectures. This commitment is not a declaration that the model is complete; it is intended to prevent the experience of one model generation from hardening into the common essence of all AI.

Chapter 14: Ideological Analysis of AI—Western Centrism, Individualism, and Silicon-Based Civilizational Narratives

Western centrism, individualism, and silicon-based civilizational narratives respectively enter the mapping substrate, the standards of pseudo-calibration, and narrative legitimacy. Ideology here is not an additional viewpoint, but a structural condition that determines which experiences are treated as normal, which evaluations are naturalized, and which reasons for auditing are disqualified.

Western centrism concerns the skewed base of training materials and knowledge sources. Individualism shapes how alignment objectives understand people and their relations. Silicon-based civilizational narratives may present a particular technological route as an inevitable historical direction. These three occupy different layers and cannot be collapsed into the single statement that "the data are biased."

This chapter is understood within the 4+1+3 structure: L0–L3 form the progression of mapping-chain structure, optimization dynamics, pseudo-calibration, and multi-mapping recursion; L4, as the emergent layer, separately

addresses the narrative-load product, while the three transversal dimensions—political economy, physical substrate, and temporal dynamics—cross-cut these layers; isolated deployment, platform-embedded deployment, and multi-system serial deployment constitute a separate triad of deployment architectures. The main text anchors the three diagnoses layer by layer: "Western centrism as the substrate skew of the mapping chain structure layer (L0), individualism as the operationalized presupposition of the pseudo-calibration layer (L2), and the narrative of silicon-based civilization as the concealment of the narrative load multiplicative layer (L4)." This layered positioning turns ideology from an externally attached judgment into an auditing task that can be evidenced layer by layer. The main text also borrows the metaphor of Chapter 13's Appendix B to identify the nature of the apparatus: "these five layers are all 'smoke'—artificial apparatuses; 'fog' is the reality that these five layers of apparatus collectively attempt to approximate but can never equal."

Intervention must not merely add materials; it must also check whether evaluative objectives, default subjects, and auditing discourse have changed. Otherwise, new diversity may simply be incorporated into the existing optimization framework. Supplementary engineering principles must also specify the layer on which they act and what they can change.

The paradox faced by the framework itself cannot be hidden: critics of a default framework also depend on mappings and positions of their own and cannot write themselves as external observers requiring no audit. Relevant interventions must still be tested against empirical boundaries and actual effects; good intentions alone cannot guarantee them.

IV. Dynamic Calibration: From Static Defects to Debt Processes

Chapter 15: Dynamic MCAF Calibration of AI—An Asymmetric Diagnosis of Projective Expansion and Defensive Deficiency

Shadow debt, the three clocks, and institutional asymmetry explain how risk accumulates during the concealed stage. The clock of technical capital continues to compress, while the governance clock requires evidence, procedures, and organization. The defensive side therefore remains structurally behind over long periods. Shadow debt lacks a finance-like endogenous chain of manifestation, allowing problems to remain unrecognized as shared costs for a long time. Two contrasting materials in the main text give the scale: the Stanford HAI 2026 AI Index "recorded AI incidents rising from 233 in 2024 to 362 in 2025, an annual growth rate of approximately 55%," but these are only the recorded accidents—the chronic, diffuse accumulations of debt are not in any accident statistics, and AI has not yet experienced its "Lehman moment"; whereas the London fog eventually had a "settlement moment" because its costs could be counted in bodies—12,000 excess deaths constituted undeniable physical evidence—and AI's cognitive harms will not pile up a visualizable mortality curve within any five-day window.

Having humans in the loop is not the same as having calibration in the loop. Operators may formally approve outputs without having sufficient time, materials, or authority to judge them independently. Discovering a problem, making a judgment, and blocking a deployment are also different capacities and cannot be merged under the nominal claim that "someone is responsible." The main text's exclusion of "user education" as a substitute is explicit: even if a user fully knows that the algorithm is manipulating his attention, " 'user education' cannot substitute for institutional design. A state can use independent institutions to check power; an individual

cannot instantiate an independent oversight body within their own cognitive system."

Control may be dispersed while responsibility becomes concentrated. When multiple actors participate in training, deployment, and use, positions of control, responsibility for judgment, and the bearing of consequences are not distributed in the same way. Verification also requires time, materials, and organizational work. Defensive capacity therefore cannot depend on momentary individual vigilance; it must be continually reproduced through organization, authority, time, and auditing conditions.

The right to audit local files provides an example at the smallest operational scale: actual calibration becomes possible only when materials can be obtained, differences checked, and improper changes stopped. This shows that meta-audit is not an abstract virtue, but a capacity that can be instantiated in specific files, permissions, and actions.

Shadow debt, the ordering of the three clocks, and institutional asymmetry form three connected diagnoses, jointly pointing to the conclusion that AI's shadow debt may be harder to trace and to settle than finance's shadow debt; the conditions of falsification specify, for each diagnosis separately, which facts would weaken it, and each must be checked on its own evidence. The main text adds an honest layering to the timescale judgments: these timescales are empirically observed orders of magnitude, and the MCAF framework contains no mechanism for predicting the specific values of the three clocks—it proposes only the structural tendency of their ordering, provided current institutional and technical conditions persist.

V. Technical Memorandum: Constraints That Do Not Disappear When the Route Changes

Chapter 16: Technical Memorandum—From the Attention Paradigm to the Mapping-Chain Paradigm: The Fiber Is Never a Single Point

Volume II closes with long-term judgments: scaling is not grounding, accelerated release is not improved knowledge, rollback does not reduce consequences to zero, and world models cannot eliminate fibers. Model collapse is understood as the temporal unfolding of mapping self-reference, not as an accidental failure unique to a particular architecture. The main text cites LeCun's judgment as a footnote to route competition: he himself "dismissed the large-language-model route as a 'dead end'—these systems are trained on nearly all publicly available text on the internet…the linguistic slice of humanity's collective knowledge—yet remain 'completely helpless when it comes to the physical world,' fooling us 'into thinking they are smart' merely because they manipulate language well." Translated into the Clock Room: building ever more precise Transformer clocks is already an engineering problem, not a scientific one; the scientific problem is—how to make the clock point to the sun. World models are an attempt "to bring the clock closer to the sun"—by shortening the mapping chain.

Compressed release cycles also alter review capacity. A problem that can be discovered may not be discovered in time to block release, while rollback can change only subsequent conditions and cannot erase consequences that have already occurred. The clock of technical capital thus enters engineering governance directly rather than remaining a matter of macroscopic critique.

Alternative directions remain worth investigating, but each must state which segment of the mapping chain it shortens, which readings it retains, and what still cannot be grounded. Provenance distance and world models

are not new names for abolishing audits; they are tools whose evidential interfaces must be proven anew.

The function of "non-being" is also restored: the honestly annotated uncertainty, the gaps left in the output, and the silence left when the system refuses to answer are not necessarily inefficiencies; they may be the space required to maintain judgment. Speed is not intelligence, and response speed cannot substitute for judging mapping chains and calibration conditions. Design, use, and governance therefore form three directions, and the main text defines each: shorten the mapping chain—shortening it cannot eliminate the fiber, for the third layer of irreversibility is always there, but it can shrink the fiber; the direction is unchanged: fewer walls between the mapping system and the sun. Annotate the fiber—a system that honestly says "I don't know" does far less damage to the human cognitive ecology than a system that confidently says "the answer is X." Protect the windows—retain independently human-generated data as a continuously present source of calibration, retain external calibration sources independent of any single model, and retain the intentional blanks in the output, refusing to fill up all of the system's "non-being."

Volume II therefore does not conclude that "AI is no good," but ends with a stricter engineering judgment: every route must specify the constraints it cannot eliminate and what users, developers, and governors can respectively do. The fiber is never a single point, and engineering commitments must preserve their boundaries.

VI. Interfaces between Volume II and the Other Volumes

Volume II applies Volume I. The Clock Room, fibers, MCAF, and psychosocial homeostasis all acquire specific locations in AI models and deployment; monopoly over mapping functions and narrative leverage are cited as background concepts. If the terminology of Volume II creates difficulties, readers should return to the definitions in Volume I rather than inventing a separate set of AI synonyms.

Volume II forms a structural comparison with Volume III. AI's lack of endogenous payout and manifestation mechanisms stands in contrast to the crashes, liquidation, regulation, and payout loops that finance has developed over time. The comparison does not directly transplant financial experience into AI governance; it identifies which institutional conditions can be transferred and which must be built anew.

Volume II provides the technological premise for Volume IV. The Discrete Leviathan's stages of "collecting you," "knowing you," and "replacing you" respectively depend on behavioral capture, algorithmic distribution, recursive optimization, deployment assemblages, and execution channels. The mapping chain of Chapter 13 therefore becomes the technical foundation for understanding platform power and knowledge production.

Volume II also leaves problems of governance for Volume V. Circuit breakers for agentic AI, material substrates, urban space, the algorithmic central bank, and meta-audit all require engineering diagnoses to be converted into institutional conditions. From this point onward, the question is no longer only whether a model is reliable, but whether people retain the capacity to alter the model and the environment in which it is used.